


Parallel Programming

Course Introduction

Phạm Trọng Nghĩa
ptnghia@fit.hcmus.edu.vn

Why need parallel?

- Many applications have demanded **more execution speed** and resources
- The rate of single-instruction stream performance scaling has decreased
 - Frequency scaling limited by power
 - **ILP** scaling tapped out
- Architects are now building faster processors by adding more execution units that run in parallel
- Software must be **written to be parallel** to see performance gains



CPU vs GPU



CPU vs GPU


CPU - Multicore

- Have **a few** cores, each core is **powerful and complex**
- Focus on **execution speed**
tốc độ thực thi



GPU – Many core

- Have **many many** cores, each core is **weak and simple**
- Focus on **throughput**



CPU vs GPU

CPU

Have **a few** cores, each core is **powerful and complex**

Focus on optimizing **latency**;
latency = an amount of time to complete a task

Example: the task is transporting a person from location A to B, the distance from A to B: 4500 km

Car: 2 people, 200 km/h
Latency = ? h $4500/200=22.5$
Throughput = ? people/h
 $2/22.5 \sim 0.09$

GPU

Have **many many** cores, each core is **weak and simple**

Focus on optimizing **throughput**;
throughput = # tasks completed in a time unit

Bus: 40 people, 50 km/h
Latency = ? h $4500/50=90$
Throughput = ? people/h
 $40/90 \sim 0.44$

CPU vs GPU

CPU

Have **a few** cores, each core is **powerful and complex**

Focus on optimizing **latency**;
latency = an amount of time to complete a task

Example: the task is transporting a person from location A to B, the distance from A to B: 4500 km

Car: 2 people, 200 km/h
Latency = **22.5** h
Throughput = **0.09** people/h

GPU

Have **many many** cores, each core is **weak and simple**

Focus on optimizing **throughput**;
throughput = # tasks completed in a time unit

Bus: 40 people, 50 km/h
Latency = **90** h
Throughput = **0.44** people/h

So, is car or bus better?

CPU vs GPU


CPU

- 24 core Intel multicore server microprocessor
- **0.33** TLOPS for double-precision and **0.66** TFLOPS for single precision




GPU - NVIDIA Tesla A100

- 108SM, 6912 CUDA cores and 432 Tensor cores
- **9.7** TFLOPS for 64-bit double-precision, **156** TFLOPS for 32-bit single-precision, and 312 TFLOPS for 16-bit half-precision



FLOPS (FLoating-point Operations Per Second)
TFLOPS (TeraFLOPS)

CPU: Latency-oriented design




- Powerful ALU
 - Reduce operation latency.
 - Increased chip area and power
- Large caches: *điện năng tiêu thụ*
 - Convert long-latency memory accesses into short-latency cache accesses
- Sophisticated control
 - Branch prediction for reduced branch latency
 - Data forwarding for reduced data latency

Reduces the execution latency of each individual thread

GPU: Throughput-oriented design

phần màu cam trong GPU là cache

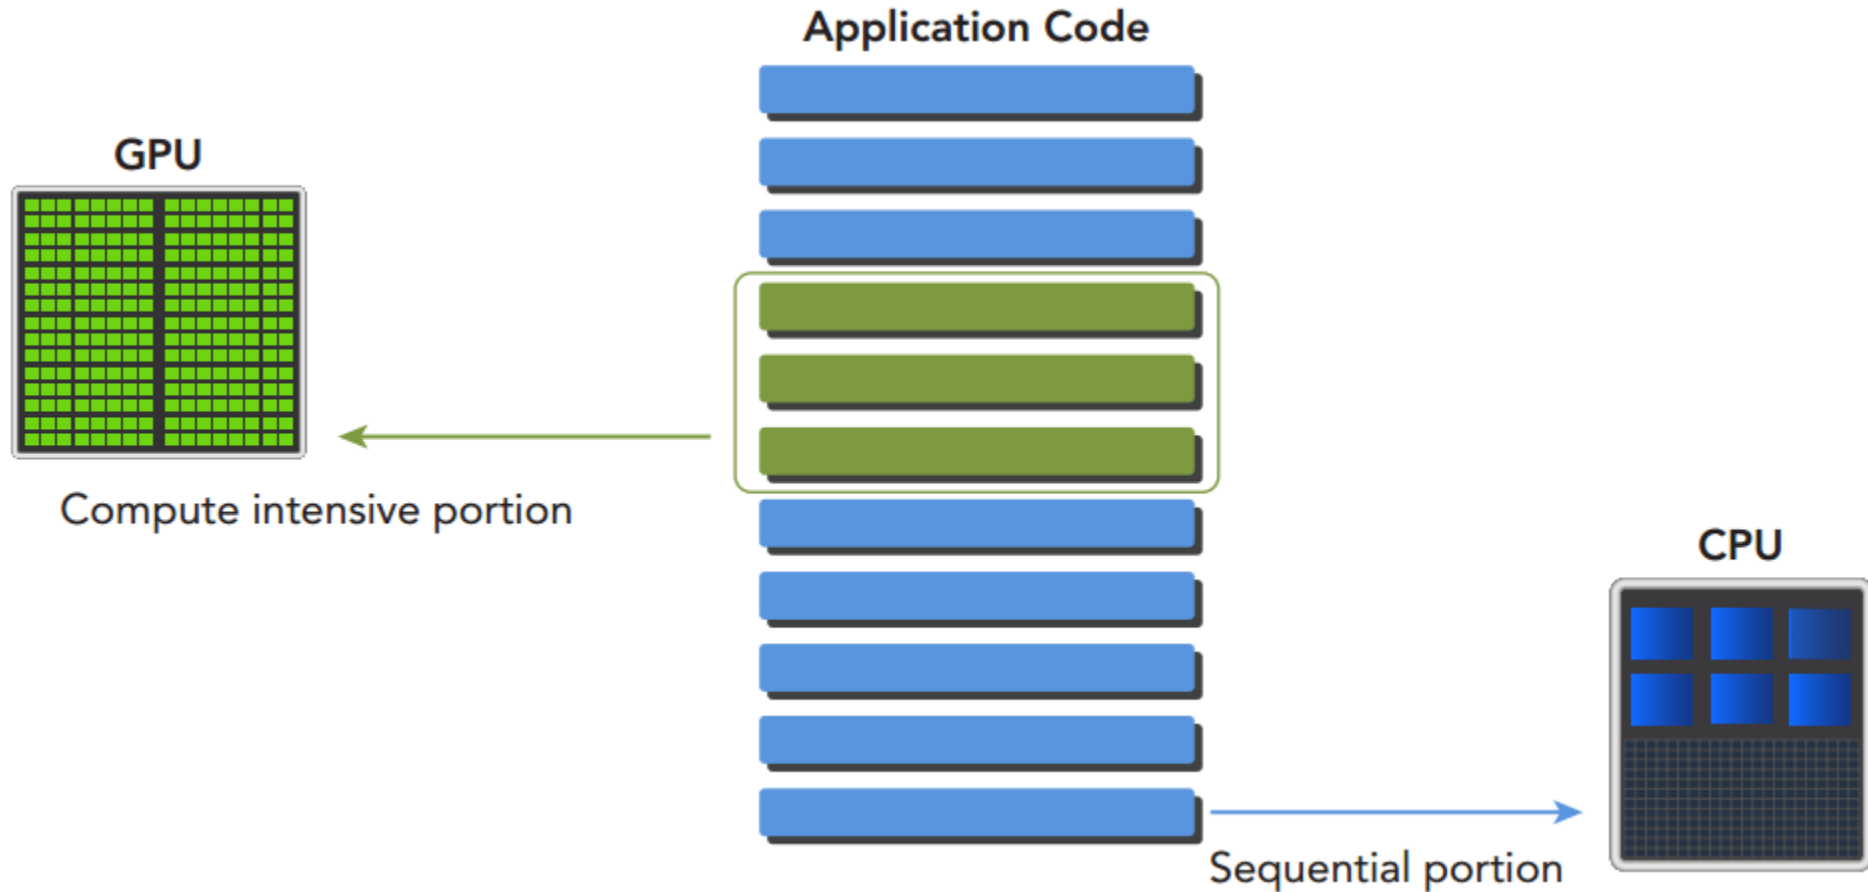


- Small caches
 - To boost memory throughput
- Simple control:
 - No branch prediction
 - No data forwarding
- Energy efficient ALUs
 - Many, long latency but heavily pipelined for high throughput
- Require massive number of threads to tolerate latencies
 - Threading logic
 - Thread size

Reduces the execution latency of each individual thread

CPU + GPU


Không thể thiếu CPU vì ta cần thao tác với OS
Dùng GPU chỉ để tăng sức mạnh tính toán



CUDA (Compute Unified Device Architecture) **C/C++** is extended-C/C++, allows us to write a program taking advantage of both CPU and GPU (NVIDIA): sequential parts will run on CPU, massively parallel parts will run on GPU

Image source: John Cheng et al. Professional CUDA C Programming. 2014


CPU + GPU




- Core area: sequential code.
 - These portions are very hard to parallelize.
 - CPUs tend to do a very good job on these portions.
 - Take up a large portion of the code, but only a small portion of the execution time
- "Peach flesh" portions:
 - Easy to parallelize.
 - Parallel programming in heterogeneous computing systems can drastically improve the speed of these applications.

Applications of parallel programming on GPU


BIOINFORMATICS



COMPUTATIONAL CHEMISTRY




COMPUTATIONAL FINANCE



COMPUTATIONAL FLUID DYNAMICS



COMPUTATIONAL STRUCTURAL MECHANICS




DATA SCIENCE



DEFENSE




ELECTRIC DESIGN AUTOMATION




IMAGING & COMPUTER VISION




MACHINE LEARNING



MEDICAL IMAGING



NUMERICAL ANALYTICS



WEATHER AND CLIMATE



Image source: <http://www.nvidia.com/object/gpu->

Challenges in parallel programming

- Question: Is parallel programming easier or hard?
- Answer:
 - **Easy**: Do not care about performance, just want it able to run.
 - **Hard**: when you want optimize, get higher performance

Challenges in parallel programming

- Challenging to design parallel algorithms with the same level of algorithmic (computational) **complexity** as that of sequential algorithms
 - Some parallel algorithms do more work than their sequential counterparts
 - không quá trực quan
 - Parallelizing often requires **non-intuitive** ways of thinking about the problem and may **require redundant work** during execution
- The execution speed of many applications is limited by **memory access** latency and/or throughput
 - Requires methods for improving memory access speed

Challenges in parallel programming

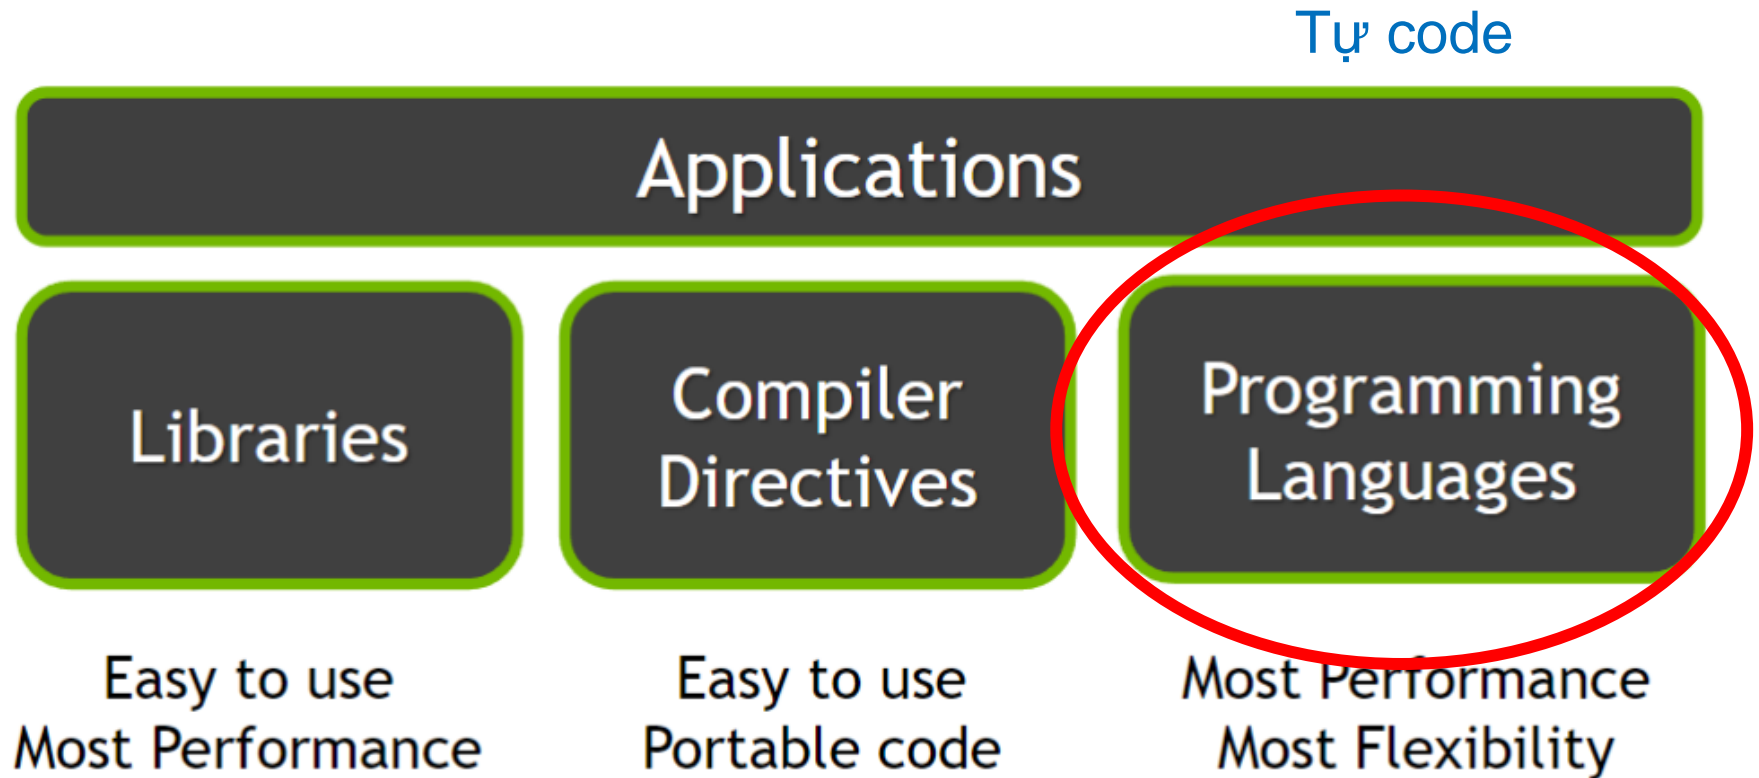
vd: thứ tự của mạng có ảnh hưởng tốc độ thuật toán quick-sort

- Execution speed of parallel programs is often more sensitive to the **input data characteristics** than is the case for their sequential counterparts
 - Unpredictable data sizes and uneven data distributions
- Require threads to collaborate with each other
 - Using synchronization operations such as barriers or atomic operations

Most of these challenges have been
addressed by researchers



3 Ways to Accelerate Applications




Libraries: Easy, High-Quality


- **Ease of use**: enables GPU acceleration without in-depth knowledge of GPU programming
- **“Drop-in”**: Many GPU-accelerated libraries follow standard APIs, thus enabling acceleration with minimal code changes
- **Quality**: Libraries offer high-quality implementations of functions encountered in a broad range of applications

NVIDIA GPU Accelerated Libraries

DEEP LEARNING





LINEAR ALGEBRA



SIGNAL, IMAGE, VIDEO



PARALLEL ALGORITHMS



<https://developer.nvidia.com/gpu-accelerated-libraries>

Compiler Directives: Easy, Portable

- **Ease of use**: Compiler takes care of details of parallelism management and data movement
- **Portable**: The code is generic, not specific to any type of hardware and can be deployed into multiple languages
- **Uncertain**: Performance of code can vary across compiler versions

Compiler Directives: OpenACC

```
// Vector_Addition.c
float * Vector_Addition
(float *restrict a, float *restrict b,
float *restrict c, int n)
{
    for(int i = 0; i < n; i ++)
    {
        c[i] = a[i] + b[i];
    }
    return c;
}

| // Vector_Addition_OpenACC.c
| float * Vector_Addition
| (float *restrict a, float *restrict b,
| float *restrict c, int n)
| {
| #pragma acc kernels loop
| copyin(a[:n], b[0:n]) copyout(c[0:n])
|     for(int i = 0; i < n; i ++)
|     {
|         c[i] = a[i] + b[i];
|     }
| }
```

<https://ulhpc-tutorials.readthedocs.io/en/latest/gpu/openacc/basics/>

Programming Languages: Most Performance and Flexible

- **Performance:** Programmer has best control of parallelism and data movement
- **Flexible:** The computation does not need to fit into a limited set of library patterns or directive types
- **Verbose:** The programmer often needs to express more details

Programming Languages: Most Performance and Flexible

Numerical analytics ▶

- MATLAB, Mathematica, LabVIEW

Python ▶

- PyCUDA, Numba

Fortran ▶

- CUDA Fortran, OpenACC

C ▶

- **CUDA C**, OpenACC

C++ ▶

- **CUDA C++**, Thrust

C# ▶

- Hybridizer

After successful completing the course, the student will be able to:

Course topics:

- ☐ Introduction to CUDA; example: vector addition, convolution, ... (3 weeks) **2 weeks**
- ☐ GPU parallel execution in CUDA; example: reduction, ... (4 weeks) **3 weeks**
- ☐ Types of GPU memories in CUDA; example: reduction, convolution, ... (3 weeks)
- ☐ Example: scan, histogram, sort (4 weeks) **3 weeks**
- ☐ Optimizing a CUDA program; additional topics in parallel programming (1 week) **0-1 week**

- ☒ Parallelize common tasks to run on GPU using CUDA
- ☒ Apply knowledge of GPU parallel execution in CUDA to speed up a CUDA program
- ☒ Apply knowledge of GPU memories in CUDA to speed up a CUDA program
- ☒ Apply the optimization process to optimize a CUDA program
- ☒ Apply teamwork skills to complete final project

Course assessment

- **Individual exercises** throughout the course: 50% of the grade
- **Group final project:** 50% of the grade, 2 students / group

Course assessment

Remember: the main goal is to **learn, truly learn**

You can discuss ideas with others as well as consult Internet sources, but **your writing and code must be your own, based on your own understanding**

If you violate this rule, you will get 0 score for the course

Advices

- In this course, we will focus on parallel programming on **GPU** (Graphics Processing Unit)
- Don't worry if you don't have GPU ;-)
- We will use Google Colab for this course.

Setup coding environment

- Where to find a machine with CUDA-enabled GPU?
 - Google Colab, it's free and ready to run CUDA programs 😊
 - Even if you have your own GPU, you should use Google Colab because teacher will use it to run and grade your programs
- Code, compile, and run:
 - Write and save code (.cu file) in your local machine by your favorite editor (with editors not recognizing .cu file automatically and not highlighting syntax with colors, the simple way is to set language/syntax as C/C++)
 - Open a notebook in [Colab](#) (you must sign in to your gmail), select “Runtime, Change runtime type” and set “Hardware accelerator” as GPU, upload .cu file
 - In a Colab cell, compile: `!nvcc file-name.cu -o run-file-name`
If we don't specify run-file-name, it will default to a.out
 - In a Colab cell, run: `!./run-file-name`
- Demo ...

RESOURCES

- Wen-Mei, W. Hwu, David B. Kirk, and Izzat El Hajj. *Programming Massively Parallel Processors: A Hands-on Approach*. Morgan Kaufmann, 2022. [Dùng chủ yếu](#)
- David B. Kirk, Wen-mei W. Hwu. *Programming Massively Parallel Processors*. Morgan Kaufmann, 2016
- Cheng John, Max Grossman, and Ty McKercher. *Professional Cuda C Programming*. John Wiley & Sons, 2014
- Lê Hoài Bắc, Vũ Thanh Hưng, Trần Trung Kiên. *Lập trình song song trên GPU*. NXB KH & KT, 2015
- NVIDIA. [Intro to Parallel Programming](#). Udacity
- NVIDIA. [CUDA Toolkit Documentation](#)

Reference

- [1] Slides from *Illinois-NVIDIA GPU Teaching Kit*
- [2] Wen-Mei, W. Hwu, David B. Kirk, and Izzat El Hajj. *Programming Massively Parallel Processors: A Hands-on Approach*. Morgan Kaufmann, 2022



THE END