

Parallel Programming

Introduction to CUDA C/C++

Part II

Phạm Trọng Nghĩa
ptnghia@fit.hcmus.edu.vn

Review: previous lecture

Structure of a simple CUDA program

- Sequential parts will run on host (CPU), massively parallel parts will run on device (GPU)
- From host, to ask device to compute in parallel:
 - Tất cả đều do host gọi
 - Host allocates memory regions on device by calling **cudaMalloc**
 - Host copies necessary data to memory regions on device by calling **cudaMemcpy**
 - Host calls **kernel function**
 - Host copies results from device by calling **cudaMemcpy**
 - Host frees memory regions on device by calling **cudaFree**

Review: previous le

Structure of a simple CUDA program

- Sequential parts will run on host (CPU), parallel parts will run on device (GPU)
- From host, to ask device to compute
 - Host allocates memory regions on device
 - Host copies necessary data to memory on device by `cudaMemcpy`
 - Host calls **kernel function**
 - Host copies results from device back to host
 - Host frees memory regions on device

- Kernel function is executed in parallel by many **threads** on device
- These threads are organized as follows:
 - The **grid** consists of **blocks** (these blocks have the same size)
 - Each block consists of threads

When host calls kernel function, host needs to specify grid size and block size

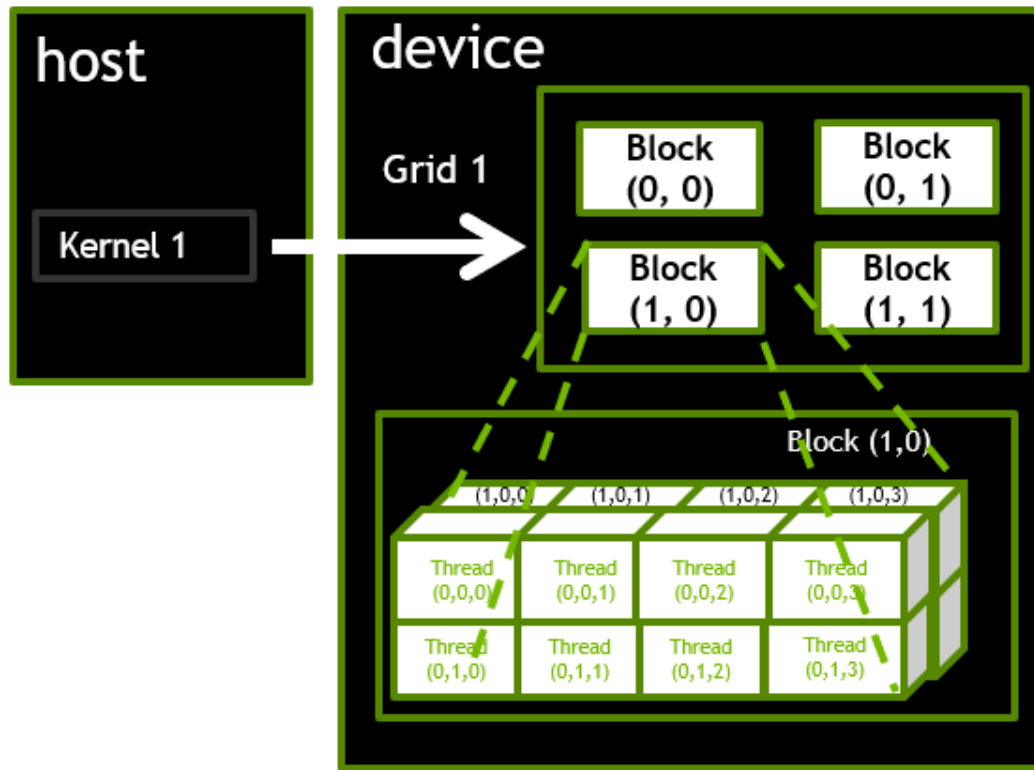
- On device, in kernel function: each thread can use built-in variables (**blockIdx**, **threadIdx**, **blockDim**, **gridDim**) to identify its data portion

Today: working with 2D data

- Last lecture:
 - Wrote a kernel function with one-dimensional grid
- This lecture:
 - More generally at *how threads are organized* and learn how threads and blocks can be used to *process multidimensional arrays*
- Multiple examples:
 - Adding 2 matrixes
 - Converting a RGB image to grayscale (intro)
 - Blurring an image (intro)

Multidimensional grid organization

- In general, a grid is a 3D array of blocks, and each block is a 3D array of threads



Multidimensional grid organization

- Typical host kernel call:

```
dim3 blockSize(16, 16, 2);  
dim3 gridSize(128, 1, 1);  
addVec<<<gridSize, blockSize >>> (...);
```

```
dim3 blockSize(16);  
dim3 gridSize(128);  
addVec<<<gridSize, blockSize >>> (...);
```

```
dim3 blockSize(16);  
dim3 gridSize(128);  
addVec<<<128, 16 >>> (...);
```

Dimension limit

Max gridDim.x: 2147483647

Max gridDim.y: 65535

Max gridDim.z: 65535

Max blockDim.x: 1024

Max blockDim.y: 1024

Max blockDim.z: 64

Max thread per block: 1024

- (512, 1, 1) ✓
- (8, 16, 4) ✓
- (32, 16, 2) ✓
- (32, 32, 2) ✗

Tổng thread:

1. 512 [v]


2. 512 [v]

3. 1024 [v]

4. 2048 [x]

Adding 2 matrixes

- **Sequential implementation:** easy 😊
- **Parallel implementation:** not difficult 😊
 - Let each thread compute one element in the output matrix
 - With 2D data, it's natural to organize threads into 2D blocks and 2D grid



blockIdx.x = 2 threadIdx.x = 3
 blockIdx.y = 1 threadIdx.y = 1

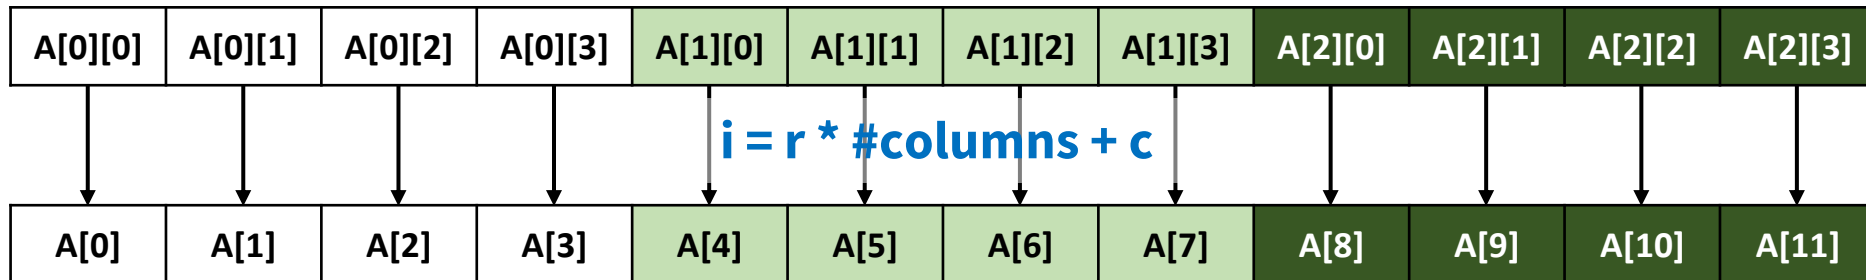

$r = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$
 $= 1 * 16 + 1$
 $= 17$
 $c = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$
 $= 2 * 16 + 3$
 $= 35$

Adding 2 matrixes

- **Sequential implementation:** easy 😊
- **Parallel implementation:** not difficult 😊
 - Let each thread compute one element in the output matrix
 - With 2D data, it's natural to organize threads into 2D blocks and 2D grid
- One final consideration before we code: how should we store a matrix?
 - Store as 2D array (a double pointer)
It will cause a lot of troubles (you can try and see ...)
 - **Store as 1D array** (a single pointer) by concatenating rows together
It will make life easier

Công việc: Trải phẳng (theo row major order)

A[0][0]	A[0][1]	A[0][2]	A[0][3]
A[1][0]	A[1][1]	A[1][2]	A[1][3]
A[2][0]	A[2][1]	A[2][2]	A[2][3]



When we implement functions operating on 2D array:

- We can **view it directly as 1D array**, but this approach doesn't work in all cases
- A more general approach is to **view it as 2D array**, compute row and column indexes, and when we need to access, convert row and column indexes to indexes in 1D array (we will use this approach)


Adding 2 matrixes: live coding

Today: working with 2D data

- Adding 2 matrixes
- Converting a RGB image to grayscale (intro)
- Blurring an image (intro)

Converting a RGB image to grayscale

Input




Output




Converting a RGB image to grayscale

Input



Output



Each pixel is represented by
3 numbers:

Each pixel is represented by
1 number:

$$\text{red, green, blue} \xrightarrow{\quad} \text{gray}$$
$$\text{gray} = 0.299 * \text{red} + 0.587 * \text{green} + 0.114 * \text{blue}$$

Each number is an unsigned
integer from 0-255

This number is an unsigned
integer from 0-255

Today: working with 2D data


- Adding 2 matrixes
- Converting a RGB image to grayscale (intro)
- Blurring an image (intro)
- CUDA functions to query device info

Blurring an image

Input



Output



Blurring an image

To blur an RGB image, we blur each channel separately

→ We just need to understand how to blur one channel

How to blur one channel?

Convolve it with a special filter to make adjacent pixels similar (details in the upcoming homework)

Today: working with 2D data

- Adding 2 matrixes
- Converting a RGB image to grayscale (intro)
- Blurring an image (intro)
- CUDA functions to query device info

Querying device info

In a CUDA program, we may want to query device info:

- For example, we want to query the max number of threads / block and the max number of blocks / grid on device and set block size and grid size appropriate for this device
→ when device is changed, code is unchanged
- Or we want to query info of all available devices in our machine and select the most powerful device to run the program
- Or we simply want to query device info and print to the screen

Querying device info

Query the number of devices

```
int devCount;  
cudaGetDeviceCount(&devCount);
```

Select a device to use (in case there are many devices), e.g. device 0

```
cudaSetDevice(0);
```

Query info of a device, e.g. device 0

```
cudaDeviceProp devProp;  
cudaGetDeviceProperties(&devProp, 0)
```

Check [here](#) for fields of `cudaDeviceProp` struct. E.g.,
`devProp.maxThreadsPerBlock` indicates the max num of threads / block

Colab device's query results (now)

Device 0: Tesla T4

Compute capability: 7.5

GMEM: 15 GB

Max number of threads per block: 1024

Max size of each dimension of a block: 1024 x 1024 x 64

Max size of each dimension of a grid: 2147483647 x 65535 x 65535

...

Reference

- [1] Wen-Mei, W. Hwu, David B. Kirk, and Izzat El Hajj.
Programming Massively Parallel Processors: A Hands-on Approach. Morgan Kaufmann, 2022



THE END