

# Parallel Programming

## Prefix sum (scan)

Phạm Trọng Nghĩa  
ptnghia@fit.hcmus.edu.vn

# Overview

---

- The “sort” task
- Sequential Radix Sort
- Parallel Radix Sort

# The “sort” task

---

in

|   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|
| 1 | 8 | 5 | 2 | 6 | 4 | 7 | 2 |
|---|---|---|---|---|---|---|---|

Stable sort

|   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|
| 1 | 2 | 2 | 4 | 5 | 6 | 7 | 8 |
|---|---|---|---|---|---|---|---|

Unstable sort

|   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|
| 1 | 2 | 2 | 4 | 5 | 6 | 7 | 8 |
|---|---|---|---|---|---|---|---|

We will focus on input array of **unsigned ints**

# Overview

---

- The “sort” task
- Sequential Radix Sort
- Parallel Radix Sort

# Sequential Radix Sort

Loop from bit b3 (least significant bit) to b1 (most significant bit):

Sort elements w.r.t. the current bit using a **stable** sort

Nhìn con số dưới dạng nhị phân, lấy bit nhanh hơn giá trị

|   | b1 | b2 | b3 |
|---|----|----|----|
| 1 | 0  | 0  | 1  |
| 0 | 0  | 0  | 0  |
| 5 | 1  | 0  | 1  |
| 2 | 0  | 1  | 0  |
| 6 | 1  | 1  | 0  |
| 4 | 1  | 0  | 0  |
| 7 | 1  | 1  | 1  |
| 2 | 0  | 1  | 0  |

|   | b1 | b2 | b3 |
|---|----|----|----|
| 0 | 0  | 0  | 0  |
| 2 | 0  | 1  | 0  |
| 6 | 1  | 1  | 0  |
| 4 | 1  | 0  | 0  |
| 2 | 0  | 1  | 0  |
| 1 | 0  | 0  | 1  |
| 5 | 1  | 0  | 1  |
| 7 | 1  | 1  | 1  |

|   | b1 | b2 | b3 |
|---|----|----|----|
| 0 | 0  | 0  | 0  |
| 4 | 1  | 0  | 0  |
| 1 | 0  | 0  | 1  |
| 5 | 1  | 0  | 1  |
| 2 | 0  | 1  | 0  |
| 6 | 1  | 1  | 0  |
| 2 | 0  | 1  | 0  |
| 7 | 1  | 1  | 1  |

|   | b1 | b2 | b3 |
|---|----|----|----|
| 0 | 0  | 0  | 0  |
| 1 | 0  | 0  | 1  |
| 2 | 0  | 1  | 0  |
| 4 | 1  | 0  | 0  |
| 5 | 1  | 0  | 1  |
| 6 | 1  | 1  | 0  |
| 7 | 1  | 1  | 1  |

unsigned int 32 bits

**DONE!**

# Sequential Radix Sort

---

- OK, Radix Sort works
- But is it efficient?

Yes, if we can make the stable sort in each loop efficient,  
e.g. work =  $O(n)$

- With unsigned int (32 bits),

Radix Sort's work  $\approx 32n = O(n)$

- It's potentially even more efficient if we process  $k > 1$  bits in each loop (and still keep the work in each loop at  $O(n)$ )

For simplicity, in this lecture, we just consider  $k=1$  bit

# Sort a binary array (corresponding to $k = 1$ bit in Radix Sort)

- Consider a binary input array:  
binIn: 0 1 1 0 1 (n elements)  
How to sort **stably** and **efficiently**? **sort giữ nguyên vị trí và complexity  $O(n)$**
- We will use Counting Sort
  - Compute the rank (the correct index in the output array) of each element (work =  $O(n)$ )  
binIn: 0 1 1 0 1  
ranks: 0 2 3 1 4  
  
Rank of binIn[i] =  
# elements < binIn[i]  
+ # elements before binIn[i] and = binIn[i]
  - Write each element to its rank in the output array (work =  $O(n)$ )  
**out[ ranks[i] ] = in[ i ]**





# Sequential Radix Sort

---

Loop from Least Significant Bit to Most Significant Bit:

Sort elements w.r.t the current bit using **Counting Sort** (stably and efficiently)

Let's implement this ...

---

Để lấy bit, ta shift phải và and 1

# Overview

---

- The “sort” task
- Sequential Radix Sort
- **Parallel Radix Sort**

# Sequential Radix Sort: parallelize?

---

Loop from Least Significant Bit to Most Significant Bit:

Sort elements w.r.t the current bit using Counting Sort

Parallelize✗

Parallelize✓

# Sort a binary array using Counting Sort: parallelize?

- Consider a binary input array:  
binIn: 0 1 1 0 1 (n elements)  
How to sort **stably** and **efficiently**?
- We will use Counting Sort
  - Compute the rank (the correct index in the output array) of each element (work =  $O(n)$ )

Parallelize ✓

- Compute # ones before each element:  
binIn: 0 1 1 0 1  
nOnesBefore: 0 0 1 2 2  
Do exclusive scan

Parallelize ✓

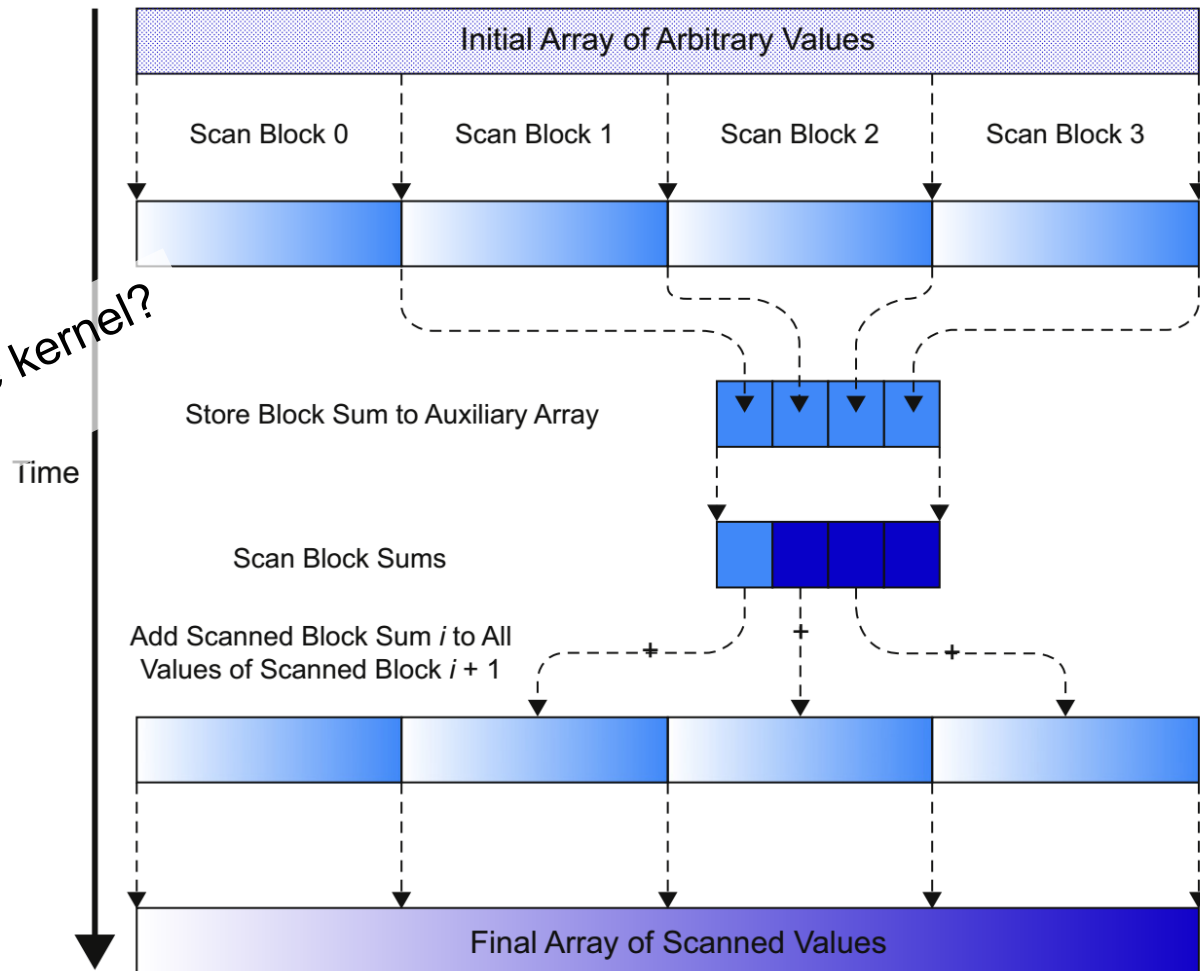
- Compute rank:  
if binIn[i] is 0: rank = i - nOnesBefore[i]  
if binIn[i] is 1: rank = nZeros + nOnesBefore[i]  
With nZeros = n - nOnesBefore[n-1] - binIn[n-1]  
binIn: 0 1 1 0 1  
ranks: 0 2 3 1 4

Parallelize ✓

- Write each element to its rank in the output array (work =  $O(n)$ )

# Remember how do we implement scan in parallel?

Can we do all 3 phases in a single kernel?  
Can we overlap these 3 phases?



# Global scan in a single kernel

---

Block with index  $bi$ :

- Scan locally
- Wait until seeing the sign indicating block  $bi-1$  has computed the sum of  $bi$  blocks ( $0 \rightarrow bi-1$ )

Get this sum, add this sum to block  $bi$ 's local sum, and turn on the sign indicating that block  $bi$  has computed the sum of  $bi+1$  blocks ( $0 \rightarrow bi$ )

(Block  $bi=0$  only needs to turn on the sign)

- Finish the rest of work: add the sum of  $bi$  blocks ( $0 \rightarrow bi-1$ ) to block  $bi$ 's local scan

(Block  $bi=0$  will not do this step)

# Global scan in a single kernel

Block with index  $bi$ :

- Scan locally
- Wait until seeing the sign indicating block  $bi-1$  has computed the sum of  $bi$  blocks ( $0 \rightarrow bi-1$ )

## A possible situation:

Blocks  $bi \rightarrow bi+N$  are assigned to available slots in SM, and wait for the result from block  $bi-1$   
Block  $bi-1$  waits for an available slot in SM

→ Deadlock 😞

**Solution:** recompute block index  $bi$ , don't tie it with  $blockIdx.x$   
Finish the rest of work and the sum of  $bi$  blocks ( $0 \rightarrow bi-1$ ) to block  $bi$ 's local scan

(Block  $bi=0$  will not do this step)

# Global scan in a single kernel

---

Block with index  $bi$ :

- Get in-order block index  $bi$
- Scan locally
- Wait until seeing the sign indicating block  $bi-1$  has computed the sum of  $bi$  blocks ( $0 \rightarrow bi-1$ )

Get this sum, add this sum to block  $bi$ 's local sum, and turn on the sign indicating that block  $bi$  has computed the sum of  $bi+1$  blocks ( $0 \rightarrow bi$ )

(Block  $bi=0$  only needs to turn on the sign)

- Finish the rest of work: add the sum of  $bi$  blocks ( $0 \rightarrow bi-1$ ) to block  $bi$ 's local scan

(Block  $bi=0$  will not do this step)



# Get in-order block index bi

- *blkCount1*: dùng để gán block ID (bid) mới. Giá trị đầu = 0
  - Block đầu tiên chạy sẽ có giá trị (bid = 0), *blkCount1++*
  - Block thứ 2 sẽ có bid = 1, *blkCount1++*
  - ...
- Chỉ thread 0 cần tính phần này vào biến share. Sau đó các thread khác lấy giá trị share này vào biến register của mình.

```
__device__ int blkCount1 = 0;
__device__ int blkCount2 = 0;
//...
if (threadIdx.x == 0)
    s_data[0] = atomicAdd(&blkCount1,1);
__syncthreads();
int bid = s_data[0];
```

# Get in-order block index bi

Hiện tại có bao nhiêu block chạy

```
__device__ int blkCount1 = 0;
__device__ int blkCount2 = 0;
//...
if (threadIdx.x == 0){
    blkSums[bid] = s_data[2 * blockDim.x - 1];
    if (bid > 0){
        while (atomicAdd(&blkCount2, 0) < bid) {}
        s_data[blockDim.x * 2] = blkSums[bid - 1];
        blkSums[bid] += s_data[blockDim.x * 2];
        __threadfence(); đảm bảo đọc/ghi đồng bộ
    }
    atomicAdd(&blkCount2, 1);
}
__syncthreads();
```

Read more:

- [Document about \\_\\_threadfence](#)
- [Document about volatile](#)

# Global scan in a single kernel

---

Block with index  $bi$ :

- Serialize between blocks*
- Get in-order block index  $bi$
  - Scan locally
  - Wait until seeing the sign indicating block  $bi-1$  has computed the sum of  $bi$  blocks ( $0 \rightarrow bi-1$ )  
Get this sum, add this sum to block  $bi$ 's local sum, and turn on the sign indicating that block  $bi$  has computed the sum of  $bi+1$  blocks ( $0 \rightarrow bi$ )  
(Block  $bi=0$  only needs to turn on the sign)
  - Finish the rest of work: add the sum of  $bi$  blocks ( $0 \rightarrow bi-1$ ) to block  $bi$ 's local scan  
(Block  $bi=0$  will not do this step)

# Inclusive scan $\overset{?}{\rightarrow}$ exclusive scan

---

# **Implement parallel Radix Sort using global scan in a single kernel**

---

The upcoming HW4 ;-)

# Radix Sort for signed ints

---

- Sign bit is MSB (Most Significant Bit)
  - MSB = 0: positive number  
Signed int = unsigned int
  - MSB = 1: negative number  
Signed int = unsigned int -  $2^{\text{\#bits-of-signed-int}}$
- If we use Radix Sort for unsigned ints, it'll be wrong
- One solution:
  - Convert signed ints to unsigned ints
  - Run Radix Sort for unsigned ints
  - Convert results back to signed ints

# Radix Sort for floats

---

- Need to understand how floats are represented
- Idea is similar to signed ints:
  - Convert floats to unsigned ints
  - Run Radix Sort for unsigned ints
  - Convert results back to floats

# Reference

---

- [1] Wen-Mei, W. Hwu, David B. Kirk, and Izzat El Hajj. *Programming Massively Parallel Processors: A Hands-on Approach*. Morgan Kaufmann, 2022
- [2] Cheng John, Max Grossman, and Ty McKercher. *Professional Cuda C Programming*. John Wiley & Sons, 2014
- [3] Illinois GPU course

<https://wiki.illinois.edu/wiki/display/ECE408/ECE408+Home>





**THE END**